

# Sistemas Distribuídos

José Orlando Pereira

Departamento de Informática  
Universidade do Minho

2018/2019



# Example: Game

- Wait for at least Min players before joining the game:

```
void joinTheGame() {  
    ready++;  
    while (ready < Min)  
        ;  
}
```

Race: `ready = ready + 1`

Busy waiting

How many problems with this code?

# Example: Game

- Wait for at least Min and at most Max players before joining the game:

```
void joinTheGame  
    ready++;  
    while (ready < Min && playing >= Max)  
        ;  
    playing++;  
}
```

Race: ready = ready + 1

Busy waiting

Race: playing = playing + 1

leaving

How many problems with this code?

# Waiting for an event

- Assuming we can ask the OS to suspend and wakeup threads, we solve busy waiting

```
void joinTheGame() {  
    ready++;  
    if (ready < Min && playing >= Max)  
        suspendMe();  
    playing++;           // playing-- when leaving  
    wakeAllOthers(); // also on leaving the game  
}
```

# Waiting for an event: 1st attempt

- Add locking:

```
void joinTheGame() {  
    l.lock();  
    ready++;  
    if (ready < Min && playing >= Max)  
        suspendMe();  
    playing++;  
    wakeAllOthers();  
    l.unlock();  
}
```



Suspended with  
lock acquired:  
**Deadlock!**

# Waiting for an event: 2nd attempt

- Unlock while suspended;

```
void joinTheGame() {
```

```
    l.lock();
```

```
    ready++;
```

```
    if (ready < Min && playing >= Max) {
```

```
        l.unlock();
```

```
        suspendMe();
```

```
        l.lock();
```

```
    }
```

```
    playing++;
```

```
    wakeAllOthers();
```

```
    l.unlock();
```

```
}
```

Unlock to sleep

Relock to update  
other variables

# Waiting for an event: 2nd attempt

- Player 1:
  - lock
  - ready++;
  - not enough ready, enter “if”
  - unlock
- Player 2:
  - lock...
  - ... acquired
  - ready++;
  - enough ready, skip “if”
  - wake suspended
  - unlock

*(between unlock and  
suspend)*

- suspended...  
forever...

**Deadlock!**

# Condition Variables

- Atomically suspends the thread and releases the lock:

```
Lock l = new ReentrantLock();
```

```
Condition c = l.newCondition();
```

```
    c.await();    // unlocks l, suspends,  
                  // and relocks l on wakeup
```

- Waking up suspended threads:
  - `c.signal();`     *// one thread*
  - `c.signalAll();`   *// all threads*

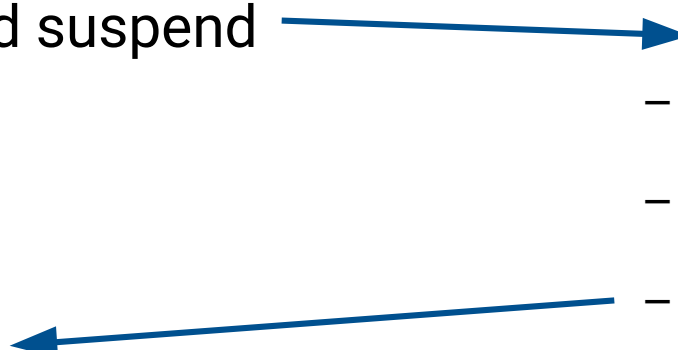


# Waiting for an event: 3rd attempt

- Atomically unlock and suspend
- Relock when waking up

```
void joinTheGame() {  
    l.lock();  
    ready++;  
    if (ready < Min && playing >= Max)  
        c.await();  
    playing++;  
    c.signalAll();  
    l.unlock();  
}
```

# Waiting for an event: 3rd attempt

- Player  $i < \text{Min}$ :
    - lock
    - ready++;
    - not enough ready, enter “if”
    - unlock and suspend
    - ...
    - ...
    - wakes up! (relock)
    - continues!
  - Player  $i = \text{Min}$ :
    - lock...
    - ... acquired
    - ready++;
    - enough ready, skip “if”
    - wake suspended
    - unlock
- 

# Waiting for an event: 3rd attempt

- Player 1:
    - lock
    - ready++;
    - not enough ready, enter “if”
    - unlock and suspend
    - ...
    - wakes up!
  - Players 2 to Max:
    - player 2 joins
    - ...
    - player Max joins
- finally gets lock!
  - playing++

**Playing > Max!**

# Waiting for an event: 4th version

- Must always use “while” loop:

```
void joinTheGame() {  
    l.lock();  
    ready++;  
    while (ready < Min && playing >= Max) {  
        c.await();  
    }  
    playing++;  
    c.signalAll();  
    l.unlock();  
}
```

Can also wake up without  
signal being called!  
("Spurious wakeup")

# Waiting for an event: General case

```
Lock l = new ReentrantLock();
```

```
Condition c = l.newCondition();
```

```
void waitForEvent() {  
    l.lock();  
    ...  
    while(!happened)  
        c.await();  
    ...  
    l.unlock();  
}  
  
void event() {  
    l.lock();  
    ... // change state  
    c.signalAll();  
    l.unlock();  
}
```

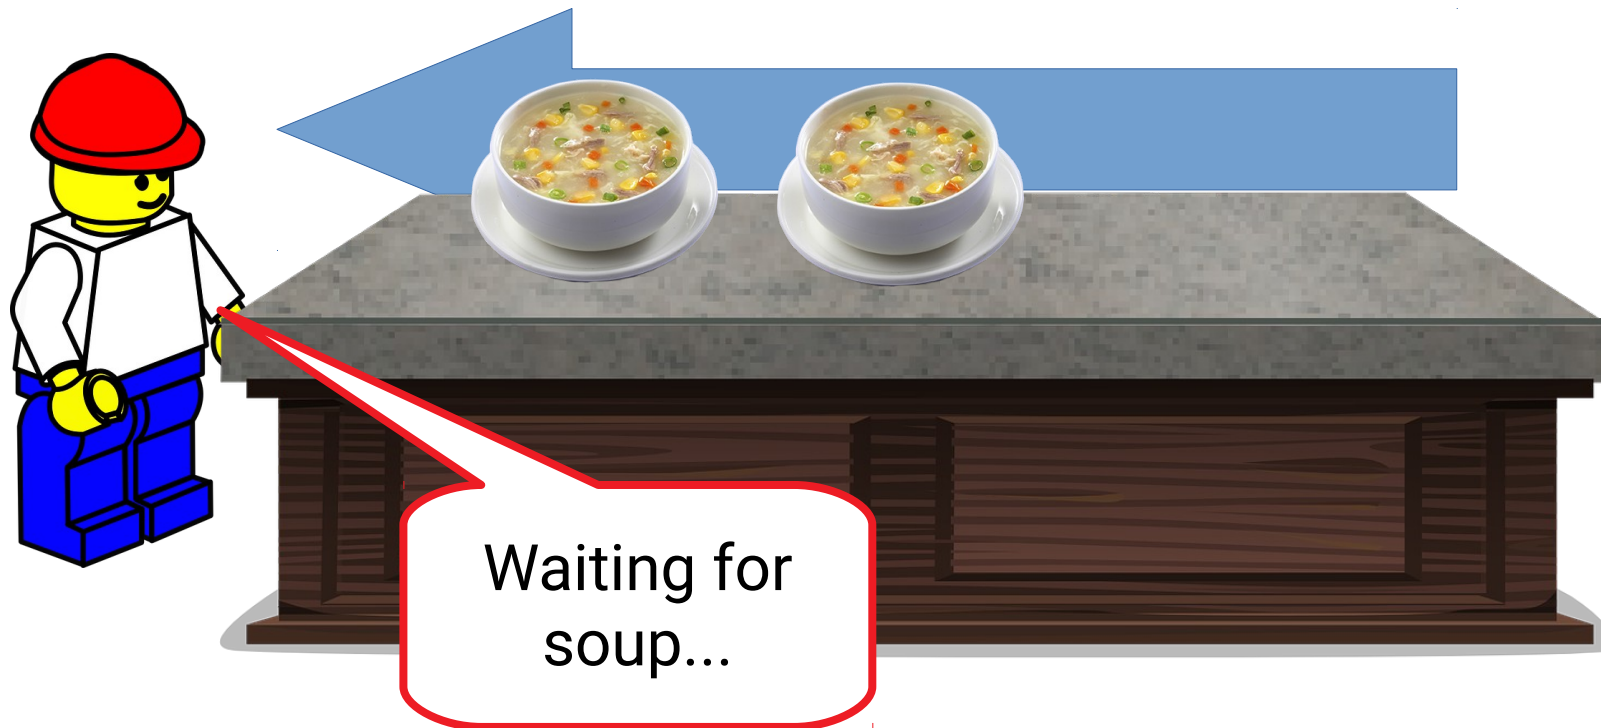
changes some value that makes the condition true

wakes up waiting threads

# Example: Soup counter



# Blocking reader



# Bounded buffer: blocking reader

```
Lock l = new ReentrantLock();  
Condition c = l.newCondition();  
Queue<Object> q = ...;
```

```
Object get() {  
    l.lock();
```

```
    while(q.isEmpty())
```

```
        c.await();
```

```
    q.remove();
```

```
    l.unlock();
```

```
}
```

```
void put(Object s) {
```

```
    l.lock();
```

```
    q.add(s);
```

```
    c.signalAll();
```

```
    l.unlock();
```

```
}
```

changes some value that  
makes the condition true

wakes up  
waiting threads

Why signal ALL  
if only one  
continues?



# Bounded buffer: blocking reader

```
Lock l = new ReentrantLock();  
Condition c = l.newCondition();  
Queue<Object> q = ...;
```

```
Object get() {  
    l.lock();
```

```
    while(q.isEmpty())
```

```
        c.await();
```

```
    q.remove();
```

```
    l.unlock();
```

```
}
```

```
void put(Object s) {
```

```
    l.lock();
```


```
    q.add(s);
```

```
    c.signal();
```

```
    l.unlock();
```

```
}
```

changes some value that  
makes the condition true



wakes up ONE  
waiting thread



# Blocking writer



# Bounded buffer: blocking writer

```
Lock l = new ReentrantLock();  
Condition c = l.newCondition();  
Queue<Object> q = ...;
```

```
Object get() {  
    l.lock();  
    q.remove();  
    c.signal();  
    l.unlock();  
}
```

```
void put(Object s) {  
    l.lock();  
    while(q.size() >= Max)  
        c.await();  
    q.add(s);  
    l.unlock();  
}
```

```
graph LR; A[q.remove()] --> B[while(q.size() >= Max)]; C[c.signal()] --> D[c.await()];
```

# Bounded buffer: blocking both

```
Lock l = new ReentrantLock();  
Condition c = l.newCondition();  
Queue<Object> q = ...;
```

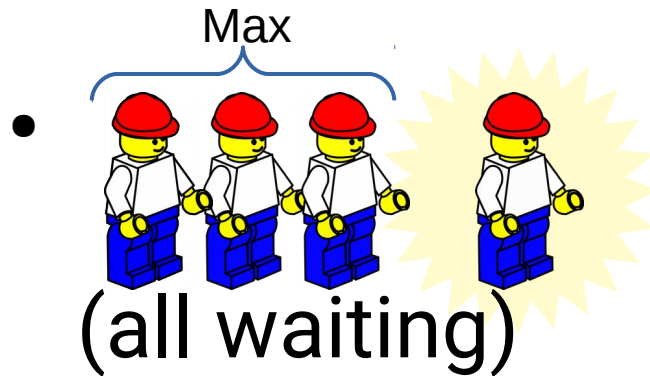
```
Object get() {  
    l.lock();  
    while(q.isEmpty())  
        c.await();  
    q.remove();  
    c.signal();  
    l.unlock();  
}
```

```
void put(Object s) {  
    l.lock();  
    while(q.size() >= Max)  
        c.await();  
    q.add(s);  
    c.signal();  
    l.unlock();  
}
```

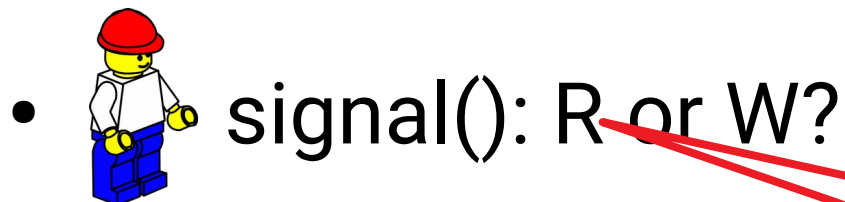
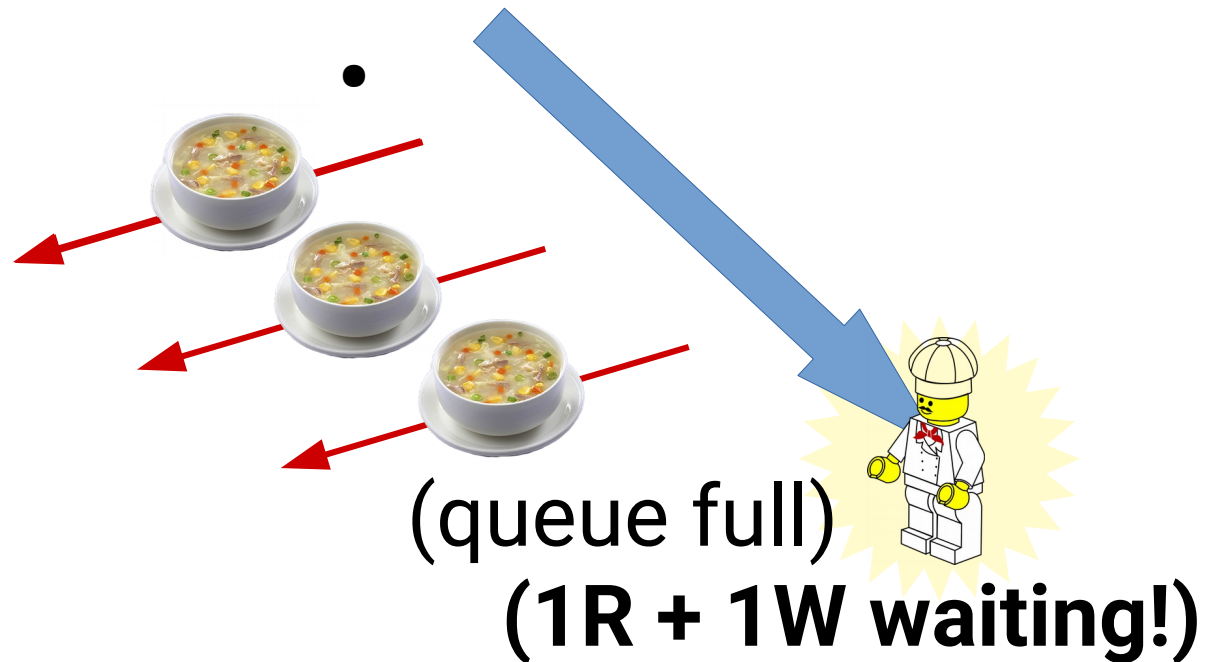
?

Is it possible to have  
readers and writers  
blocked at the same time?

# Bounded buffer



(3 waiting)  
(2 waiting)  
(1 waiting)



**Deadlock**

# Bounded buffer

- This is a general problem with different conditions and the same condition variable
- Solution: Use `signalAll()`
  - This is the only solution with Java “synchronized” monitors (with `notifyAll()`)
  - Inefficient (“thundering herd”)



# Bounded buffer: 2 condition variables

- Lock l = new ReentrantLock();  
Condition notEmpty = l.newCondition();  
Condition notFull = l.newCondition();  
Queue<Object> q = ...;

```
Object get() {  
    l.lock();  
    while(q.isEmpty())  
        notEmpty.await();  
    q.remove();  
    notFull.signal();  
    l.unlock();  
}
```

```
void put(Object s) {  
    l.lock();  
    while(q.size() >= Max)  
        notFull.await();  
    q.add(s);  
    notEmpty.signal();  
    l.unlock();  
}
```

# Conclusions

- Use `await()` always within a “while” loop
  - Races
  - Spurious wakeups
- Minimizing wakeups improves performance and scale
  - Use one condition variable for each condition to avoid `signalAll()`
- Blocking bounded buffer is an important building block in concurrent programming



# j.u.c Conditions vs Monitors

- class C {

There is a hidden "condition" in each object used by "wait()/notify()"

```
synchronized public void m1() {  
    while(...) wait();  
}
```

```
synchronized public void m2() {  
    notify();  
}
```

```
}
```

- class C {

```
    private Lock l =  
        new ReentrantLock();
```

```
    private Condition c =  
    - l.newCondition();
```

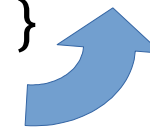
```
    public void m() {  
        try { l.lock();  
            while(...) c.await()  
        } finally { l.unlock(); }  
    }
```

```
    public void m() {  
        try { l.lock();  
            c.signal();  
        } finally { l.unlock(); }  
    }
```

```
}
```



Equivalent code  
(aproximately...)



# j.u.c. Conditions vs Monitors

- Main differences:
  - One implicit condition for each lock  
vs.  
Many j.u.c. conditions for the same lock
    - Avoids signalAll()
  - Threads waiting for a condition wakeup in any order  
vs.  
Threads waiting for a j.u.c. Condition wakeup in  
FIFO order