

# Sistemas Distribuídos

José Orlando Pereira

Departamento de Informática  
Universidade do Minho

2018/2019




# Application Level Multicast

- Send to multiple destinations: group
- Reliability:
  - Deliver all messages sent
  - vs
  - All deliver the same messages (agreement)


# General approach

- Buffer and retransmit until acknowledged




# Acknowledgments

- Not scalable to large number of destinations due to “ack implosion”:




# Group acknowledgment

- Not scalable to high throughput due to large buffer requirements and slow recovery:




# Structured multicast

- Avoids “ack implosion” but increases GC latency:




# Crybaby

- A single slow receiver delays GC, regardless of structure:



# Local acknowledgment

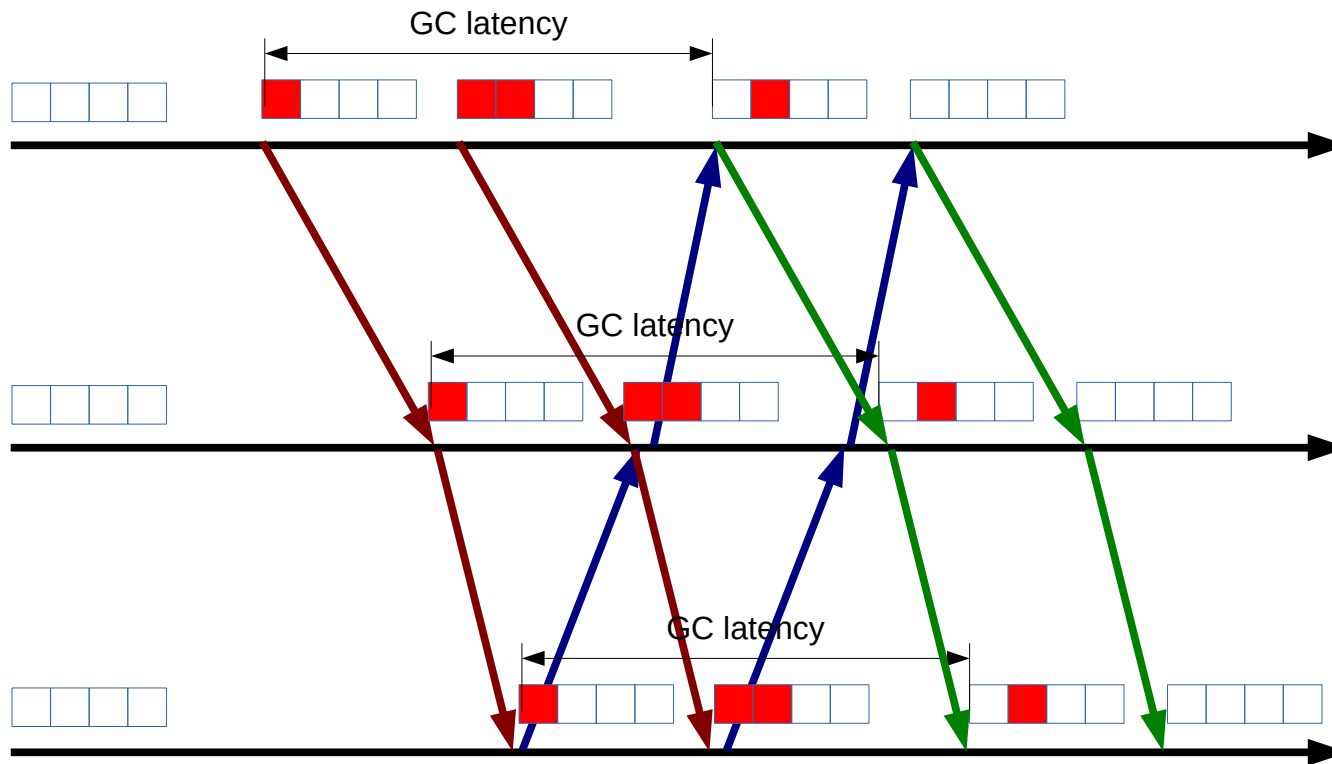
- Local acknowledgment decreases GC latency by using additional buffering:





# Multicast with Agreement


- Two rounds of acknowledgment required to avoid “ $O(n^2)$  ack implosion”:




# Gossip

- Simple protocol to multicast a message:
  - Select a small subset of random targets
  - Forward message only to those targets
  - Discard message
- Upon receiving a new message, act as the sender

# Gossip



# Gossip




# Gossip and Epidemics

- Similarity with epidemics:
  - Sender = contagious = spreads rumor
  - Receiver = infected = knows rumor
  - Ignores duplicated = dead = old news...
- Interesting parameters:
  - $n$  – size of the population
  - $f$  – number of targets

# Fanout vs Reliability


- Infected rate  $\pi$ :
  - $\pi = 1 - \exp(-\pi f)$
  - Independent of  $n$ !
- Probability of atomic infection  $p$ :
  - $f = \log(n) + c$ ,  $p = \exp(-\exp(-c))$
  - Depends on  $n$ !
- Duration of epidemic when infecting the entire population order of  $\log(n)$

# Fanout vs Reliability



(50000 destinations)


# Redundancy





# Peer Sampling and Overlays


- Need random sample of  $f$  peers from the set of targets
- Current sample defines implicit overlay network:



- Cannot assume knowledge of all targets to draw from:
  - Potentially very large list
  - How to update it dynamically (churn)?


# Desirable Overlay Properties

- Connectivity!
  - Convergence to  $f+c$
  - Low variance for load balancing
- Low diameter
  - Latency
- Uniform sample
  - Avoid clustering




# Proactive Overlay Maintenance

- Keep a list of known targets
- Periodically shuffle with known peers:
  - Evict some local targets at random
  - Adopt some remote targets at random



# Reactive Overlay Maintenance

- To join a group, use a contact point
- Initiate  $f+c$  random walks in the overlay from the contact point
- Upon arrival, insert new member in local target set



# Fault Tolerance

- Pros:
  - Tolerates faults by increasing parameter  $f$ :
    - Packet loss
    - Dead processes
  - Immune to the crybaby, since there is no feedback
- Cons:
  - Assumes independent faults

# Performance

- Pros:
  - No tree setup/repair overhead
  - Perfect load balancing
- Cons:
  - Redundancy